

Assignment 2 - Spark Dataframes

***Note*:** All the dataset files were stored in the same folder as this notebook.

Table of Contents:

- [Dependencies](#)
- [Instructions on running the code](#)
- [Start a Spark Session](#)
- [Question 1](#)
- [Question 2](#)
- [Question 3](#)
- [Question 4](#)
- [Question 5](#)
- [Question 6](#)
- [Question 7](#)

Dependencies


JDK 8

Download link: <https://www.oracle.com/java/technologies/javase/javase8-archive-downloads.html>

You will need to create an account and sign-in in order to download it.

Press start and search for "Edit environment variables for your account"

Click new and add 'JAVA_HOME' with the path for your java jdk folder e.g. C:\Program Files\Java\jdk1.8.0_202 and add 'JRE_HOME' with the path for your java jre folder e.g. C:\Program Files\Java\jre1.8.0_361



PySpark

PySpark installation using PyPI is as follows:

```
pip install pyspark
```

```
pip install pyarrow
```

If you want to install extra dependencies for a specific component, you can install it as below:

```
# Spark SQL
```

```
pip install pyspark[sql]
```

```
# pandas API on Spark
```

```
pip install pyspark[pandas_on_spark] plotly # to plot your data, you can install plotly together
```

More info about about installing PySpark: [PySpark Installation Guide](#)

Haversine

```
pip install haversine
```

More info about about Haversine: <https://pypi.org/project/haversine/>

Instructions on running the code:

After you [start a spark session](#), you can run each cell in the notebook independently.

Start a Spark Session

[Table of Contents](#)

```
In [1]: import os
import pyspark
import sys

os.environ['PYSPARK_PYTHON'] = sys.executable
os.environ['PYSPARK_DRIVER_PYTHON'] = sys.executable
conf = pyspark.SparkConf()
sc = pyspark.SparkContext(conf=conf)
spark = pyspark.sql.SparkSession(sc)
spark
```

Out[1]: **SparkSession - in-memory**

SparkContext

[Spark UI](#)

Version	v3.3.2
Master	local[*]
AppName	pyspark-shell

1. 15 Points

Datafile: BreadBasket_DMS.csv

Solve: What is the most popular (most sold) between the 8:00AM and 8:59AM for each day?

Example output (not actual solution)

2016-10-30, Pastry

2016-10-31, Coffee

:

:

Approach:

1. Import BreadBasket_DMS.csv into a dataframe and filter out rows with `NONE` in the `Item` column
2. Extract dates in `YYYY-MM-DD` format from the `Date` column and times in `hh:mm:ss` format from the `Time` column
3. Filter the data by `Time` in the range of `08:00:00` and `08:59:00` inclusive
4. Group the data by `Date` and `Item`, aggregate by the `count` and, sort by `Date` and `count`
5. Group the data by `Date` and return the `last Item` and `last count`
6. Inner join with the ungrouped data over `Date` and `count` to include ties

Table of Contents

```
In [2]: # 1. Import BreadBasket_DMS.csv into a dataframe and filter out rows with `NONE`
# in the `Item` column
from pyspark.sql.functions import col
BreadBasket_DMS = spark.read\
.option("header", True)\
.option("inferSchema", True)\
.csv("BreadBasket_DMS.csv")\
.filter(col("Item") != "NONE")

# 2. Extract dates in `YYYY-MM-DD` format from the `Date` column and times in
# `hh:mm:ss` format from the `Time` column
from pyspark.sql.functions import to_date, date_format
BreadBasket_DMS = BreadBasket_DMS\
.withColumn("Date", to_date(col("Date"), "YYYY-MM-DD"))\
.withColumn("Time", date_format(col("Time"), "hh:mm:ss"))

# 3. Filter the data by `Time` in the range of `08:00:00` and `08:59:00` inclusive
q1 = BreadBasket_DMS
q1 = q1.filter((col("Time") <= "08:59:00") & (col("Time") >= "08:00:00"))
```

```

# 4. Group the data by `Date` and `Item`, aggregate by the `count` and, sort by
# `Date` and `count`
q1 = q1.groupBy("Date", "Item").count().sort("Date", "count")
q1.persist()

# 5. Group the data by `Date` and return the Last `Item` and Last `count`
from pyspark.sql.functions import last
q1_no_ties = q1.groupBy("Date").agg(
    last("Item").alias("Most Popular Item"),
    last("count").alias("Total Count"))

# 6. Inner join with the ungrouped data over `Date` and `count` to include ties
JoinExpression = (
    q1["Date"] == q1_no_ties["Date"]) & (q1["count"] == q1_no_ties["`Total Count`"])
q1_with_ties = q1\
    .join(q1_no_ties, JoinExpression, "left_semi")\
    .select(
        "Date",
        col("Item").alias("Most Popular Items"),
        col("count").alias("Total Counts"))

# Display results
print("1. List of the most popular (most sold) items between the 8:00 AM and " +
      "8:59 AM for each day and their total transactions that day:")
print("\nWITH TIES:")
q1_with_ties.show()
print("WITHOUT TIES:")
q1_no_ties.show()

```

1. List of the most popular (most sold) items between the 8:00 AM and 8:59 AM for each day and their total transactions that day:

WITH TIES:

Date	Most Popular Items	Total Counts
2016-10-31	Coffee	2
2016-10-31	Bread	2
2016-11-01	Tea	3
2016-11-02	Coffee	8
2016-11-03	Coffee	4
2016-11-04	Bread	2
2016-11-04	Coffee	2
2016-11-05	Bread	6
2016-11-07	Pastry	1
2016-11-07	Coffee	1
2016-11-08	Bread	1
2016-11-08	Pastry	1
2016-11-08	Coffee	1
2016-11-09	Pastry	1
2016-11-09	Bread	1
2016-11-09	Coffee	1
2016-11-10	Coffee	2
2016-11-11	Bread	6
2016-11-12	Coffee	1
2016-11-12	Medialuna	1

only showing top 20 rows

WITHOUT TIES:

Date	Most Popular Item	Total Count
2016-10-31	Bread	2
2016-11-01	Tea	3
2016-11-02	Coffee	8
2016-11-03	Coffee	4
2016-11-04	Coffee	2
2016-11-05	Bread	6
2016-11-07	Coffee	1
2016-11-08	Coffee	1
2016-11-09	Coffee	1
2016-11-10	Coffee	2
2016-11-11	Bread	6
2016-11-12	Medialuna	1
2016-11-14	Coffee	2
2016-11-15	Keeping It Local	1
2016-11-16	Bread	1
2016-11-17	Siblings	2
2016-11-18	Coffee	6
2016-11-19	Bread	3
2016-11-21	Coffee	2
2016-11-22	Medialuna	1

only showing top 20 rows

2. 15 Points

Datafile: BreadBasket_DMS.csv

Solve: What is the most common item bought along with "Brownie"? (items bought in the same transaction)

Assumptions:

We will assume that we will count each time an item was bought with "Brownie". If an item was bought more than once in the same transaction we will count each time that item was bought in that transaction.

Approach:

1. Import BreadBasket_DMS.csv into a dataframe and filter out rows with `NONE` in the `Item` column
2. Extract dates in `YYYY-MM-DD` format from the `Date` column and times in `hh:mm:ss` format from the `Time` column
3. Create 2 lists of transactions: one with and one without "Brownie" in the `Item` column by using `filter()`
4. Use a `left_semi` join where the `Transaction` is the same only keeping items bought with "Brownie"
5. Group the data by `Item`, aggregate by `count()` of the data and, sort the data by `count` in descending order

Table of Contents

```
In [3]: # 1. Import BreadBasket_DMS.csv into a dataframe and filter out rows with `NONE`
# in the `Item` column
from pyspark.sql.functions import col
BreadBasket_DMS = spark.read\
    .option("header", True)\
    .option("inferSchema", True)\
    .csv("BreadBasket_DMS.csv")\
    .filter(col("Item") != "NONE")

# 2. Extract dates in `YYYY-MM-DD` format from the `Date` column and times in
# `hh:mm:ss` format from the `Time` column
from pyspark.sql.functions import to_date, date_format
BreadBasket_DMS = BreadBasket_DMS\
    .withColumn("Date", to_date(col("Date"), "YYYY-MM-DD"))\
    .withColumn("Time", date_format(col("Time"), "hh:mm:ss"))
```

```

# 3. Create 2 Lists of transactions: one with and one without "Brownie" in the
# `Item` column by using `filter()`
BrownieTrnsctns = BreadBasket_DMS.filter(col("Item") == "Brownie")
OtherTrnsctns = BreadBasket_DMS.filter(col("Item") != "Brownie")

# 4. Use a `left_semi` join where the `Transaction` is the same only keeping items
# bought with "Brownie"
JoinExpression = BrownieTrnsctns["Transaction"] == OtherTrnsctns["Transaction"]
BoughtWithBrownie = OtherTrnsctns\
.join(BrownieTrnsctns, JoinExpression, "left_semi")\
.sort("Transaction")

# 5. Group the data by `Item`, aggregate by `count()` of the data and, sort the
# data by `count` in descending order
from pyspark.sql.functions import desc
BoughtWithBrownie = BoughtWithBrownie.groupBy("Item").count().sort(desc("count"))

# Display results
BoughtWithBrownie.persist()
print("2. The most common item bought with along with "Brownie" is " +
      f"{BoughtWithBrownie.first()[0]}" which was purchased " +
      f"{BoughtWithBrownie.first()[1]} times")
print("List of most common items bought along with "Brownie" sorted by count:")
BoughtWithBrownie.show()

```

2. The most common item bought with along with "Brownie" is "Coffee" which was purchased 237 times

List of most common items bought along with "Brownie" sorted by count:

Item	count
Coffee	237
Bread	115
Tea	71
Cake	43
Hot chocolate	42
Alfajores	27
Sandwich	27
Cookies	26
Juice	24
Pastry	23
Medialuna	19
Muffin	18
Soup	15
Scone	12
Farm House	11
Coke	11
Truffles	11
Mineral water	9
Toast	7
Hearty & Seasonal	6

only showing top 20 rows

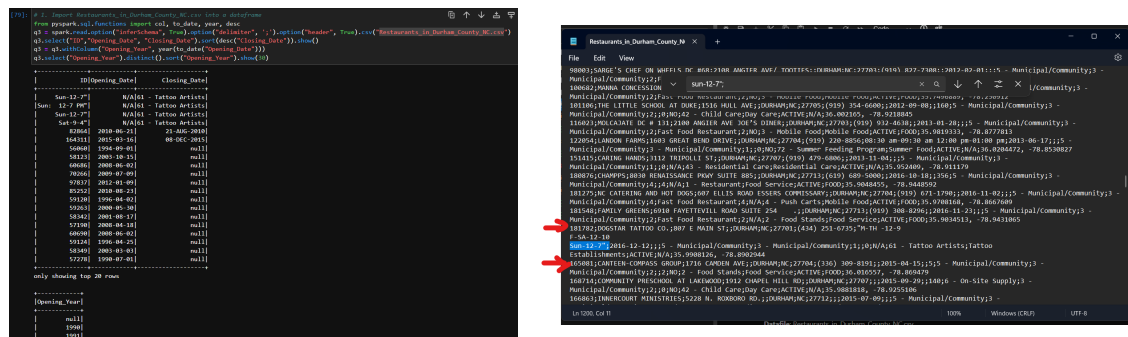
3. 10 Points

Datafile: Restaurants_in_Durham_County_NC.csv

NOTE* This file is colon delimited (not comma) *****

Solve: How many years are represented in this dataset?

Inspecting the Restaurants_in_Durham_County_NC.csv file, I noticed that there several line breaks in the file that cause PySpark to think that this is a new row. Here is an example of this:



To fix this, I used `.option("multiline", True)`

Source: <https://stackoverflow.com/questions/48800046/escape-new-line-character-in-spark-csv-read>

Approach:

1. Import `Restaurants_in_Durham_County_NC` into a dataframe
2. Extract the year in "yyyy" format from the `Opening_Date` column into a new `Opening_Year` column
3. Group by `Opening_Year` and aggregate over `count()`
4. Extract the `min`, `max` and and total cout of years represented in the data

Extracting single value from DataFrame: <https://www.geeksforgeeks.org/pyspark-extracting-single-value-from-dataframe/>

Table of Contents

```
In [4]: # 1. Import Restaurants_in_Durham_County_NC.csv into a dataframe
# Source for reading csv files with multiline:
# stackoverflow.com/questions/48800046/escape-new-Line-character-in-spark-csv-read
Restaurants = spark.read\
.option("inferSchema", True)\
.option("delimiter", ';')\
.option("header", True)\
```



```

.option("multiline", True)\
.csv("Restaurants_in_Durham_County_NC.csv")

# 2. Extract the year in "yyyy" format from the `Opening_Date` column into a new
# `Opening_Year` column
from pyspark.sql.functions import col, to_date, year
Restaurants = Restaurants.withColumn("Opening_Year", year(to_date("Opening_Date")))

# 3. Group by `Opening_Year` and aggregate over `count()`
Years = Restaurants.groupBy("Opening_Year").count()
Years.persist()

# 4. Extract the `min`, `max` and and total cout of years represented in the data
# Extracting single value from DataFrame:
# https://www.geeksforgeeks.org/pyspark-extracting-single-value-from-dataframe/
from pyspark.sql.functions import min, max, first
Year_min = Years.select(min("Opening_Year")).first()[0]
Year_max = Years.select(max("Opening_Year")).first()[0]
Year_count = Years.count()

# Display the results
print(f"3. There are {Year_count} years are represented in this dataset. " +
      f"They span from {Year_min} to {Year_max}")
Years.show(Year_count)

```

3. There are 28 years are represented in this dataset. They span from 1990 to 2017

Opening_Year	count
1990	60
2003	64
2007	68
2015	209
2006	49
2013	156
1997	23
1994	192
2014	195
2004	55
1991	31
1996	62
1998	61
2012	125
2009	81
2016	253
1995	116
2001	48
1992	12
2005	69
2000	29
2010	94
2011	125
2008	151
2017	32
1999	47
2002	42
1993	14

4. 20 Points

Dataset: Restaurants_in_Durham_County_NC.csv

Solve: Show the type and count of restaurant opened during the 90's (1990-1999 inclusive).

Note: type="Rpt_Area_Desc"

Example (not the actual result):

"Swimming Pools", 13

"Tattoo Establishment", 2

:

Approach:

1. Import Restaurants_in_Durham_County_NC.csv into a dataframe
2. Extract the year in "yyyy" format from the Opening_Date column into a new Opening_Year column
3. Filter the Opening_Year column between 1990 and 1999 inclusive
4. Group by the column Rpt_Area_Desc and aggregate by count

Table of Contents

```
In [5]: # 1. Import Restaurants_in_Durham_County_NC.csv into a dataframe
# Source for reading csv files with multiline:
# stackoverflow.com/questions/48800046/escape-new-line-character-in-spark-csv-read
Restaurants = spark.read\
.option("inferSchema", True)\
.option("delimiter", ';')\
.option("header", True)\
.option("multiline", True)\
.csv("Restaurants_in_Durham_County_NC.csv")

# 2. Extract the year in "yyyy" format from the `Opening_Date` column into a new
# `Opening_Year` column
from pyspark.sql.functions import col, to_date, year
Restaurants = Restaurants.withColumn("Opening_Year", year(to_date("Opening_Date")))

# 3. Filter the `Opening_Year` column between 1990 and 1999 inclusive
q4 = Restaurants\
.filter((col("Opening_Year") <= 1999) & (col("Opening_Year") >= 1990))

# 4. Group by the column `Rpt_Area_Desc` and aggregate by `count`
q4 = q4.groupBy("Rpt_Area_Desc").count()

# Display the results
print("4. List of the type and count of restaurants opened during the 90's " +
      "(1990-1999 inclusive):")
q4.show()
```

4. List of the type and count of restaurants opened during the 90's (1990-1999 inclusive):

```
+-----+-----+
| Rpt_Area_Desc | count |
+-----+-----+
| Bed&Breakfast Home | 3 |
| Summer Camps | 1 |
| Institutions | 16 |
| Local Confinement | 2 |
| School Buildings | 56 |
| Swimming Pools | 256 |
| Day Care | 58 |
| Bed&Breakfast Inn | 1 |
| Lodging | 21 |
| Food Service | 204 |
+-----+-----+
```

5. 25 Points

Dataset: populationbycountry19802010millions.csv

Solve: For region, compute the ***percentage change*** in population, year over year. Note the year 1980 will not have a preceding year. For each year, display the region with the top population ***decrease***.

Example (not actual results):

1981, North America, -2%

1982, Aruba, -7%...

Assumptions:

The region with the top population ***decrease*** has the lowest ***percentage change***.

Approach:

1. Import populationbycountry19802010millions.csv in to a dataframe and rename `_c0` to `country`
2. Initialize an empty array called `q5_rows` to store the rows for the result
3. Loop through all the years
4. Find `Percentage Change` and sort by sort it in ascending order of `Percentage Change`
5. Add a `Year` column and store the year corresponding to the change
6. Append the first row which is the top population decrease or the lowest percentage change into `q5_rows`
7. Create a schema for a dataframe and store the rows from `q5_rows`

Table of Contents

```
In [6]: # 1. Import populationbycountry19802010millions.csv in to a dataframe and rename
# `_c0` to `country`
population = spark.read\
.option("header", True)\
.option("InferSchema", True)\
.csv("populationbycountry19802010millions.csv")\
.withColumnRenamed("_c0", "country")

# 2. Initialize an empty array called `q5_rows` to store the rows for the result
q5_rows = []

# 3. Loop through all the years
from pyspark.sql.functions import col, first, lit
for year in range(1980, 2010):
    # 4. Find `Percentage Change` and sort by sort it in ascending order of
```

```

# `Percentage Change`
_q5 = population.withColumn(
    "PercChng", (col(str(year+1)) - col(str(year)))*100/col(str(year)))\
.dropna().\
sort("PercChng")
# 5. Add a `Year` column and store the year corresponding to the change
_q5 = _q5.withColumn("Year", lit(year+1))
# 6. Append the first row which is the top population decrease or the lowest
# percentage change into `q5_rows`
q5_rows.append(
    _q5.select(
        "Year",
        "Country",
        col("PercChng").alias("Percentage Change"))
    .first())

# 7. Create a schema for a dataframe and store the rows from `q5_rows`
from pyspark.sql.types import StructField, StructType, StringType, IntegerType, DoubleType
q5_schema = StructType([
    StructField("Year", IntegerType(), True),
    StructField("Country", StringType(), True),
    StructField("Percentage Change", DoubleType(), True)])
q5 = spark.createDataFrame(q5_rows, q5_schema)

# Display the results
print("5. List of the regions with the top population decrease for 1980 to 2010:")
q5.show(30)

```

5. List of the regions with the top population decrease for 1980 to 2010:

Year	Country	Percentage Change
1981	Afghanistan	-9.106330931425992
1982	Afghanistan	-8.017227257036874
1983	Antigua and Barbuda	-3.5141890898397343
1984	Antigua and Barbuda	-1.7525144772935055
1985	Cook Islands	-1.4092446448703508
1986	Netherlands Antilles	-24.58781655279631
1987	Saint Helena	-21.299638989169676
1988	Mozambique	-2.883631837516533
1989	Somalia	-2.1964965331028314
1990	Liberia	-12.816300240117076
1991	Kuwait	-55.4531619095637
1992	Somalia	-5.387440289087448
1993	Bosnia and Herzeg...	-7.072116841830721
1994	Rwanda	-14.363511428676736
1995	Rwanda	-15.871881307134093
1996	Montserrat	-22.590068159688407
1997	Montserrat	-25.157232704402517
1998	Montserrat	-43.193277310924366
1999	Cook Islands	-2.9919447640966608
2000	Cook Islands	-3.2621589561091247
2001	Cook Islands	-3.55610055180871
2002	Cook Islands	-3.6872218690400547
2003	Montserrat	-6.652806652806653
2004	Djibouti	-4.830771012478634
2005	Montserrat	-8.669354838709674
2006	Nauru	-4.39560439560439
2007	Nauru	-4.702194357366778
2008	Cook Islands	-3.3096926713948007
2009	Cook Islands	-3.259983700081494
2010	Cook Islands	-3.2013479359730423

6. 15 Points

Dataset: romeo-juliet-pg1777.txt

Solve: Do **word count** in pyspark. Ignore punctuation, and normalize to lower case. Accept only the characters in this set: **[0-9a-zA-Z]**

Discussion

For the purpose of this assignment, all characters not in this set: [a-z, A-Z, 0-9] were replaced with spaces or " " using the following `regex_string` :

```
regex_string = r"[^a-zA-Z0-9]"
```

However, there are are issues with this and one such issue is discussed below:

Hyphenated and Apostrophe Words:

The regex used would mistreat characters with hyphens or apostrophes such as `Don't` or `Mother-in-law` into `["don", "t"]` and `["mother", "in", "law"]`.

Proposed Alternative Regex:

A more appropriate regex to include hyphenated and appstrophe words might be given by the two options below. (source: <https://stackoverflow.com/questions/27715581/>):

```
re.findall(r"(?!'.*')\b[\w'-]+\b", line.lower())
```

or

```
re.findall(r"[A-Za-z0-9]+(?:['-][A-Za-z0-9]+)*", line.lower())
```

However, this would also include possessive words such as `professor's` which might be undesired.

Approach:

1. Import romeo-juliet-pg1777.txt in to a dataframe
2. Replace words not in `[0-9a-zA-Z]` with " "
3. Transform the words in the dataframe to lower case
4. `split` words in line into an array
5. `explode` the arrays of words in line
6. Group by word and aggregate over count

Table of Contents

```
In [7]: # 1. Import romeo-juliet-pg1777.txt in to a dataframe
RomeoJuliet = spark.read.text("romeo-juliet-pg1777.txt")
WordCount = RomeoJuliet

# 2. Replace words not in `[0-9a-zA-Z]` with " "
from pyspark.sql.functions import regexp_replace
regex_string = r"^[a-zA-Z0-9]"
WordCount = WordCount.select(
    regexp_replace("value", regex_string, " ").alias("LINE"))

# 3. Transform the words in the dataframe to lower case
from pyspark.sql.functions import lower
WordCount = WordCount.select(lower("LINE").alias("line"))

# 4. `split` words in line into an array
from pyspark.sql.functions import split
WordCount = WordCount.select(split("line", " ").alias("words_in_line"))

# 5. `explode` the arrays of words in line
from pyspark.sql.functions import explode
WordCount = WordCount.select(explode("words_in_line").alias("word"))
```

```
# 6. Group by word and aggregate over count
WordCount = WordCount.groupBy("word").count()

# Display the results
print("6. Word Count of romeo-juliet-pg1777.txt")
WordCount.show()
```

6. Word Count of romeo-juliet-pg1777.txt

```
+-----+-----+
|      word|count|
+-----+-----+
|    those|   17|
|  carnegie|   10|
|    some|   58|
|    chor|    2|
|    art|   55|
|   still|   15|
|  nourish|    1|
|    cures|    1|
|solemnity|    3|
|    feign|    1|
|   imagin|    1|
|consortest|    1|
|   pitcher|    1|
|    earl|    1|
|    hope|    4|
|   shroud|    3|
|   unfirm|    1|
|   embrace|    1|
|    often|    4|
|  received|    3|
+-----+-----+
```

only showing top 20 rows

7. Extra credit – 30 points

Datasets:

Restaurants_in_Durham_County_NC.csv

durham-nc-foreclosure-2006-2016.json

Solve: For each restaurant ('Restaurants_in_Durham_County_NC.csv' with "status"="ACTIVE" and "rpt_area_desc"="Food Service", show the number of foreclosures ('durham-nc-foreclosure-2006-2016') within a radius of **1 mile** of the restaurant's coordinates.

Note: Use any assumption for the shape of Earth...

Or you can use the Haversine distance. <https://pypi.org/project/haversine/>

Note: UDF, or user defined functions, is part of next week's lecture

Approach:

1. Import Restaurants_in_Durham_County_NC.csv into a dataframe named Restaurants
2. filter the restaurants for where "status" == "ACTIVE" and "rpt_area_desc" == "Food Service"
3. Extract the latitudes and longitudes values from the geolocation column, cast them as double type, and drop nulls
4. Import the fields.geocode column from durham-nc-foreclosure-2006-2016.json into a dataframe and drop nulls
5. Define a function that returns True if 2 points are within specified radius and make it a udf called distance_udf
6. Join Restaraunts with Foreclosure using distance_udf as the join condition

Table of Contents

```
In [8]: # 1. Import Restaurants_in_Durham_County_NC.csv into a dataframe named `Restaurants`
Restaurants = spark.read \
    .option("inferSchema", True) \
    .option("delimiter", ';') \
    .option("header", True) \
    .option("multiline", True) \
    .csv("Restaurants_in_Durham_County_NC.csv")

# 2. `filter` the restaurants for where `status` == "ACTIVE"
# and `rpt_area_desc` == "Food Service"
from pyspark.sql.functions import col
Restaurants = Restaurants.filter(
    (col("status") == "ACTIVE") &
    (col("rpt_area_desc") == "Food Service"))

# 3. Extract the latitudes and longitudes values from the `geolocation` column,
# cast them as double type, and drop nulls
from pyspark.sql.functions import split
Restaurants = Restaurants.select(
    "Premise_Name",
    split("geolocation", ", ")[0].cast("double").alias("lat"),
    split("geolocation", ", ")[1].cast("double").alias("lng"))\
    .dropna()

# 4. Import the `fields.geocode` column from durham-nc-foreclosure-2006-2016.json
# into a dataframe and drop nulls
Foreclosure = spark.read\
    .json("durham-nc-foreclosure-2006-2016.json")\
    .select("fields.geocode")\
    .dropna()

# 5. Define a function that returns `True` if 2 points are within specified radius
# and make it a udf called `distance_udf`
from pyspark.sql.functions import udf
```

```

from pyspark.sql.types import BooleanType
from haversine import haversine, Unit

def distance(lat, lng, geocode, unit='mi', radius=1):
    return haversine((lat, lng), geocode, unit=unit) <= radius

distance_udf = udf(distance, BooleanType())

# 6. Join `Restaraunts` with `Foreclosure` using `distance_udf` as the join condition
q7 = Restaurants.join(
    Foreclosure,
    distance_udf(Restaurants["lat"], Restaurants["lng"], Foreclosure["geocode"]))

```

```

In [9]: q7.persist()
        q7.count()

```

Out[9]: 63593

```

In [10]: # Display Results
print("7. The number of foreclosures within a radius of 1 mile of the active, food"
      + " service restaurants' coordinates:")
q7.groupBy("Premise_Name").count().show()

```

7. The number of foreclosures within a radius of 1 mile of the active, food service restaurants' coordinates:

```

+-----+-----+
|      Premise_Name|count|
+-----+-----+
|    DPAC 3RD FLOOR|  267|
|W G PEARSON SCHOO...|  199|
|  COMPARE FOODS DELI|   65|
|GSK COMMERCIAL OP...|    1|
|  MCDONALD'S 35265|   17|
|DUKE UNIVERSITY W...|    6|
|  BLU SEAFOOD AND BAR|   31|
|      JADE BUFFET|   33|
|HARRIS TEETER 172...|   10|
|      GRILL 46|   17|
|  EL DORADO'S #6|    5|
|  THE BLUE NOTE GRILL|  111|
|  CHICK-FIL-A EXPRESS|    6|
|HARRIS TEETER 224...|    9|
|      MATEO|  198|
|NZINGA'S BREAKFAS...|  346|
|      SPROUT CAFE|   57|
|      WENDY'S 6316|   39|
|LA SUPERIOR MEAT ...|   39|
|    PAGE ROAD GRILL|    2|
+-----+-----+

```

only showing top 20 rows